

CHARLES BENELLO

+1(312) 989-9950 | cmbenello@uchicago.edu | [linkedin.com/in/charles-benello](https://www.linkedin.com/in/charles-benello) | [cmbenello.github.io](https://github.com/cmbenello)

EDUCATION

The University of Chicago

Chicago, IL

Master of Science in Financial Mathematics

Expected Dec 2026

- Wayne Booth Graduate Student Teaching Prize (2026) - UChicago's top graduate teaching award
- Coursework: Distributed Systems, ML in SWE, Agentic AI, Reinforcement Learning, Computing for Finance

The University of Chicago | GPA: 3.75/4.0

Chicago, IL

Bachelor of Science in Mathematics and Computer Science (Systems), Minor in Korean

Jun 2025

- Coursework: Database Systems, Operating Systems, Algorithms, Concurrency Control, Computer Architecture, Networks

EXPERIENCE

Database Research Assistant

Apr 2023 – Present

ChiData Research Lab, University of Chicago

Chicago, IL

- First-authored VLDB 2026 paper (CrocSort) by designing a parallel external merge sort in Rust that completes under tight memory budgets where PostgreSQL, DuckDB, and ClickHouse abort, jointly planning memory and phase-specific thread counts.
- Co-authored CIDR 2025 paper on resource-adaptive query execution with paged memory management, enabling query operators to adapt to available memory at runtime under contention.
- Sole-authored HPTS 2026 paper on a lightweight prediction model for sort configuration on cloud and disaggregated hardware, transferring calibration across machines within 3–5% of optimal without brute-force sweeps.

GPU Systems Researcher

Jun – Sep 2025

EPFL (DIAS Lab, Prof. Anastasia Ailamaki)

Lausanne, Switzerland

- Engineered a macroblock-based GPU decompression strategy in CUDA/C++ that cut out-of-memory query times by 50–200% on large analytical workloads.
- Built cost-based simulation models predicting query performance across workload complexities, enabling adaptive decompression strategy selection at plan time.
- Designed join patterns for TPC-DS benchmarks with resource contention models across I/O, decompression, and compute pipelines.
- Benchmarked GPU compression trade-offs against CPU baselines, shaping a new cost-based research direction with a forthcoming publication.

Software Engineering Intern

Feb 2026 – Present

Alter Domus

Chicago, IL

- Lifted document extraction accuracy from 50% to 92–97% by architecting a multi-stage pipeline with schema-enforced structured output and cross-reference verification for 100+ page financial filings.
- Built an LLM-as-a-Judge evaluation framework with automated benchmarks for faithfulness, numerical accuracy, and completeness; cut LLM inference costs by optimizing context handling.

Quantitative Researcher

Mar 2026 – Present

PayPal (UChicago Project Lab)

Chicago, IL

- Built a self-supervised transaction foundation model (TabBERT-style hierarchical 2-level transformer, ~5M params) pretrained on 14.5M unlabeled transactions, reaching AUC 0.974 / F1 0.822 on downstream fraud detection at 0.12% positive rate.
- Ran 87+ ablation experiments on UChicago's Midway3 SLURM cluster across architecture, tokenization, time encoding, masking rate, and sequence length; implemented per-field masking and SPRM from PANTHER (NeurIPS 2025) and PRAGMA (Revolut 2026), with per-field masking alone yielding +1.9% AUC and +16.8% F1.

PROJECTS

GPU CrocSort | CUDA, C++ | github.com/cmbenello/gpu-research

2025–Present

- Ported CrocSort to GPU with a novel shared-memory K-way merge tree ($\log_2 K$ levels of 2-way merge-path inside shared memory), bridging the prior tradeoff between high merge fan-in and full intra-block parallelism. 8-way merge cuts HBM traffic $3\times$ vs. 2-way merge-path on 19,531 runs (10 GB vs. 30 GB on 1 GB data); targets 85–90% of A100/H100 HBM peak with adaptive cost-based config selection across radix vs. K-way merge.

SQL Query Executor | Rust | github.com/cmbenello/query_exec

2024

- Built a SQL query executor in Rust supporting all 22 TPC-H benchmark queries, including ORDER BY, CTEs, mark joins, semi-joins, subqueries, IN-lists, casts, and complex filter predicates.

Relational Database Engine | Rust | github.com/cmbenello/CrustyDB

2023

- Implemented a multi-threaded relational database from scratch: B+ tree indexing, concurrent buffer pool with page-level locking, heap storage, and a SQL query execution engine.

TECHNICAL SKILLS

Languages: C/C++, Rust, CUDA, Python, SQL, Go, TypeScript, JavaScript

Tools: PostgreSQL, DuckDB, Databricks, Docker, Git, Linux, PyTorch, Next.js, FastAPI

Concepts: Database Internals, Query Execution & Optimization, External Sort, Distributed Systems, GPU Programming, Concurrency, Cost-Based Optimization, Performance Engineering

PUBLICATIONS

[1] R. Otaki*, C. Benello*, et al. "CrocSort: Resource-Efficient, Skew-Resilient Parallel External Merge Sort." *VLDB 2026*. (*Equal contribution)

[2] C. Benello. "Getting the Most Out of External Sorting in Cloud and Disaggregated Environments." *HPTS 2026*.

[3] R. Otaki, J.H. Chang, C. Benello, A. Elmore, G. Graefe. "Resource-Adaptive Query Execution with Paged Memory Management." *CIDR 2025*.

ADDITIONAL

Languages: Korean, Italian (Intermediate); French (Basic) | **Interests:** Running (Chicago Marathon '22, '23, '24), Climbing, Weightlifting